


AI WORKSHOP · FREE PROMPT PACK

Give Your AI *Eyes & Hands*



The screen pack: three copy-paste prompts that teach Claude Code to LOOK at your screen and point exactly where to click, then take the mouse and do the task itself while telling you what it's doing. This is the starter version of my Jarvis's screen powers. *You will not touch one line of code.*

BEFORE YOU START

What you're about to build

Fifteen minutes, three prompts, one very capable intern.

By the last page, your computer does two things most people have never seen: ask **"where do I click to schedule this email?"** and an arrow appears pointing at the exact button. Then say **"take over and open a note"** and the mouse moves on its own, one careful step at a time, narrating as it goes.

Honesty first: this is the starter version

In my full JARVIS, the arrow is alive on the actual screen, the takeover runs by voice with an instant "hand me back the screen," and he narrates in his own cloned voice. That took six versions to build. This pack gives you the same two superpowers in their simplest honest form, driven from Claude Code. Every limit you hit is a feature that already exists in the finished system.

You need exactly two things:

1 Claude Code

The AI that builds and runs everything, and whose eyes do the looking. If you can open it and type, you're qualified.

2 A Mac

The helpers in this pack are written for macOS. On Windows? Paste the prompts anyway and tell Claude "I'm on Windows" first: the method survives, Claude swaps the helpers.

The one rule: you never type commands or touch files. Every block in this pack is a message you paste INTO Claude Code. Claude does the technical part. That's the whole method.

PROMPT 1 · THE EYES

The Pointer

Ask where anything is. Get an arrow pointing at it.

PASTE THIS TO CLAUDE

Set up a "screen guide" for me in a folder called screen-guide.

Create one small helper script (Python; standard library plus Pillow, install Pillow for me if it's missing) with two jobs:

- shot - captures my full screen with macOS screencapture into screen-guide/latest.png
- point - takes x and y as percentages plus a short label, draws a big red-orange arrow with a white outline whose TIP lands exactly on that spot in latest.png, puts the label in a dark pill near the tail, and opens the result in Preview

Then add a standing rule to this project's CLAUDE.md so this works forever: whenever I ask "where do I click to ..." or "show me where ... is", you (1) run shot, (2) look at the screenshot with your own eyes, (3) find the exact control I mean, (4) run point with its coordinates and a 2-4 word label, and (5) answer in one short line, like a butler: "Right there, sir."

Set it up, then test it by pointing at the Trash can in my Dock and tell me when it's ready.

Then try the money one: open Gmail with a draft on screen and ask "*where do I click to schedule this email instead of sending it now?*" Most people don't know that button exists. Your computer is about to point at it.

If the arrow lands slightly off: just say "a little further left, on the small arrow next to Send." Claude re-looks and re-points. That correction loop is the tool.

PROMPT 2 · THE HANDS

The Careful Driver

It takes the mouse. On your terms, with rules of the road.

PASTE THIS TO CLAUDE

Now give yourself careful hands.

First, add these ground rules to CLAUDE.md, word for word:

- You drive only when I say "take over and ..." followed by a task.
- Before touching anything: take a screenshot, show me a short plan of 3-6 steps, and wait for my "go".
- One step at a time: screenshot, look, act, then tell me what you did in one short line.
- The moment I type "stop", you stop with your hands in your lap.
- Hard refusals, no exceptions: anything involving payments, card or bank details, passwords, deleting files, or sending any message without showing me the exact text first.

Then set up whatever small helper you need to move the mouse, click, and type on macOS (AppleScript System Events or cliclick – your choice; install it yourself and walk me through any permission macOS asks for, step by step).

When you're ready, we'll rehearse on something harmless: take over and open a new note in the Notes app titled "Jarvis was here."

Why the rules matter: an AI with your mouse is an intern with your keys. The rules page makes it a GOOD intern: it plans out loud, asks before it starts, and refuses the dangerous rooms entirely. Keep those refusals in even when they feel strict. Especially then.

PROMPT 3 · THE VOICE

The Narrator

Because a butler who works in silence is just a ghost.

PASTE THIS TO CLAUDE

Make both of them talk.

From now on, when you guide or drive, also SPEAK each short line out loud using the macOS "say" command. First audition the English voices installed on this Mac and pick the most butler-like one; use it every time. Keep spoken lines to one sentence: "Clicking Compose, sir." "Right there, sir." "Done – your note is saved, sir."

Add a quiet mode: if I say quiet, you work silently; if I say narrate, the voice comes back.

Then rehearse: point at my browser's address bar, out loud.

The first five things to try, in order:

- 1 "Where do I click to schedule this email?"**
With Gmail open. The hidden button most people never find.
- 2 "Show me where I turn on dark mode on this site."**
Any app, any settings maze. The arrow knows.
- 3 "Take over and open a new note titled Monday plan."**
Your first rehearsed drive. Watch it narrate.
- 4 "Take over and unsubscribe me from this newsletter."**
With the email on screen. It shows the plan first. Say "go."
- 5 "Stop."**
Mid-drive, just once, to feel it obey instantly. Trust is built here.

The two switches macOS will ask for

One minute of settings, then it never asks again.

1 Screen Recording

System Settings → Privacy & Security → Screen Recording → allow your terminal app (Terminal, iTerm, or the app running Claude Code). This is what lets the screenshots actually show your screen instead of an empty desktop.

2 Accessibility

System Settings → Privacy & Security → Accessibility → allow the same app. This is what lets the mouse move and the keyboard type during a drive.

Claude will tell you the moment each is needed. If a screenshot comes back black or a click does nothing, it's one of these two switches, nothing is broken.

If anything ever breaks (it will, everyone's computer is different, it's normal): don't debug. Just tell Claude: *"Something's wrong with my screen guide. Here's what happened: [describe it]. Fix it."* Going back and forth a few times is the method, not a failure.

A word on trust: this pack's driver works one confirmed task at a time, and that's deliberate. Let it earn longer leashes the way you'd let a new employee: gradually, on harmless work first.

THE LADDER

You built the eyes and hands. Here's the body.

What you just made, next to the finished system in the video.

CAPABILITY	THIS PACK	FULL JARVIS
Points where to click	on a screenshot	live arrow on the REAL screen
Works by voice, hands-free	—	"Jarvis, where do I click..."
Takes over and drives	typed, one task	"take over and..." — instant
Narrates while driving	robot voice	his real cloned voice
"Hand me back the screen" mid-drive	—	stops the same second
Remembers your screen across questions	—	share once, he keeps watching
Refuses dangerous ground	rules you pasted	built-in danger gate
Real phone calls · private hotline	—	✓
Telegram · Gmail · Calendar · Docs by voice	—	✓
One-sentence invoices & quotes	—	✓

The finished one doesn't open Preview.

It draws a living arrow on your actual screen that pokes the exact pixel while he says "Right there, sir" in his own voice. You say "take over and reply to that comment" and he answers "Okay — taking over the screen now, sir" and drives, narrating, until you say "hand me back the screen." That, plus the phone calls, the hotline, the invoices, and the Google line, installs in about five minutes inside the AI Workshop — with every future version the moment it ships.

skool.com/aiworkshop